



Customer Churn

Springboard Capstone 2 Milestone Report 1

Brock Nosbisch

10.01.2019

Problem Statement

My client is a natural gas company. Being able to predict when a customer will turn off is crucial to the natural gas business. By knowing when a customer will turn off, you can forecast future turn off work and also allow sales reps to reach out to those customers to try to get them to stay on.

The Dataset

The dataset used in this project comes from both internal data and external data sources. The data will be merged inside our internal data warehouse and then the dataset was manually downloaded and imported into the Data Wrangling notebook.

The original data was sourced from multiple applications and external sources. It was combined in a data warehouse and then exported with the below columns. The imported data set contains 670,000 records with 93 columns.

Int64Index: 670890 entries, 98617454736 to 119996415468 Data columns (total 93 columns):				
CIS_DIVISION	667276	non-null	object	
ACCOUNT_ID	670890	non-null	int64	
CUSTOMER_CLASS_CODE	670890	non-null	object	
CUSTOMER_CLASS_DESCRIPTION	670890	non-null	object	
PERSON_ID	670890	non-null	int64	
SA_START_DATE	670890	non-null	object	
SA_START_YEAR	670890	non-null	int64	
SA_START_YEAR_MONTH	670890	non-null	int64	
SA_END_DATE	670890	non-null	object	
SA_END_YEAR	670890	non-null	int64	
SA_END_YEAR_MONTH	670890	non-null	int64	
SA_STATUS_FLAG	670890	non-null	int64	
SA_TYPE_CODE	670890	non-null	object	
RATE_CLASS_CODE	670451	non-null	object	
RATE_CLASS_DESCRIPTION	670451	non-null	object	
PREMISE_ID	670890	non-null	int64	
CITY	670890	non-null	object	
STATE	670890	non-null	object	
POSTAL	670890	non-null	int64	
BILL_CYCLE_CODE	665926	non-null	object	
BILL_CYCLE_DESCRIPTION	665926	non-null	object	
SERVICE_TYPE_CODE	670890	non-null	object	
PREMISE_TYPE_CODE	670890	non-null	object	
TREND_AREA_CODE	670890	non-null	object	
OFFICE_LOCATION	670847	non-null	object	
OFFICE_LOCATION_DESCRIPTION	670847	non-null	object	
DNP_STARTED_FLAG	670890	non-null	object	
DNP_STOPPED_FLAG	670890	non-null	object	
INTERNAL_CREDIT_RATING	670890	non-null	int64	
PREMISE_LEVEL_12_MTH_DNP_FLAG	670890	non-null	object	
ACCOUNT_LEVEL_12_MTH_DNP_FLAG	670890	non-null	object	
PERSON_RCD_18_MTHS_PLEDGE	670890	non-null	object	
PAYMENTS_IN_LAST_18_MONTHS	670890	non-null	int64	
BILLS_IN_LAST_18_MONTHS	670890	non-null	int64	
PAY_SEGS_IN_LAST_18_MONTHS	670890	non-null	int64	
BILL_SEGS_IN_LAST_18_MONTHS	670890	non-null	int64	
ARREARS_CURRENT_AMOUNT	670890	non-null	float64	
ARREARS_PAYOFF_AMOUNT	670890	non-null	float64	
TOTAL_CURRENT_AMOUNT	670890	non-null	float64	
TOTAL_PAYOFF_AMOUNT	670890	non-null	float64	
MOST_RECENT_PAYMENT_DATE	569710	non-null	object	
LATE_PAYMENT_COUNT	670890	non-null	int64	
SEASONAL_PRIOR_1_YR_FLAG	670890	non-null	object	
SEASONAL_PRIOR_2_YR_FLAG	670890	non-null	object	
SEASONAL_PRIOR_3_YR_FLAG	670890	non-null	object	
SA_START_DEGREE_DAY	670890	non-null	int64	
SA_START_AVG_TEMP	670890	non-null	int64	
SA_END_DEGREE_DAY	670890	non-null	int64	
SA_END_AVG_TEMP	670890	non-null	int64	
STOP_2013	670890	non-null	int64	
STOP_2014	670890	non-null	int64	
STOP_2015	670890	non-null	int64	
STOP_2016	670890	non-null	int64	
STOP_2017	670890	non-null	int64	
STOP_2018	670890	non-null	int64	
STOP_2019	670890	non-null	int64	
START_2013	670890	non-null	int64	
START_2014	670890	non-null	int64	
START_2015	670890	non-null	int64	
START_2016	670890	non-null	int64	
START_2017	670890	non-null	int64	
START_2018	670890	non-null	int64	
START_2019	670890	non-null	int64	
PLEDGE_DATE_2013	0	non-null	float64	
PLEDGE_FLAG_2013	670890	non-null	int64	
PLEDGE_DATE_2014	0	non-null	float64	
PLEDGE_FLAG_2014	670890	non-null	int64	
PLEDGE_DATE_2015	4426	non-null	object	
PLEDGE_FLAG_2015	670890	non-null	int64	
PLEDGE_DATE_2016	5920	non-null	object	
PLEDGE_FLAG_2016	670890	non-null	int64	
PLEDGE_DATE_2017	5729	non-null	object	
PLEDGE_FLAG_2017	670890	non-null	int64	
PLEDGE_DATE_2018	3141	non-null	object	
PLEDGE_FLAG_2018	670890	non-null	int64	
PLEDGE_DATE_2019	0	non-null	float64	
PLEDGE_FLAG_2019	670890	non-null	int64	
USAGE_IN_LAST_18_MONTHS	425368	non-null	float64	
PERSON_MAX_SA_START_DATE	668553	non-null	object	
PERSON_MIN_SA_START_DATE	668553	non-null	object	
PERSON_MAX_SA_END_DATE	480944	non-null	object	
PERSON_MIN_SA_END_DATE	480944	non-null	object	
PREMISE_MAX_SA_START_DATE	669656	non-null	object	
PREMISE_MIN_SA_START_DATE	669656	non-null	object	
PREMISE_MAX_SA_END_DATE	458841	non-null	object	
PREMISE_MIN_SA_END_DATE	458841	non-null	object	
PREMISE_PRIOR_STOP_DATE	283942	non-null	object	
PERSON_PRIOR_STOP_DATE	176492	non-null	object	
PREMISE_DAYS_INACTIVE_BEFORE	670890	non-null	int64	
PERSON_DAYS_INACTIVE_BEFORE	670890	non-null	int64	
PREMISE_DAYS_ACTIVE_BEFORE	670890	non-null	int64	
PERSON_DAYS_ACTIVE_BEFORE	670890	non-null	int64	
ACTIVE_DIFF_FROM_20190301	380148	non-null	float64	

Data Wrangling

As mentioned, most of the data manipulation and joins were done in SQL prior to loading the data into the Python notebook.

The dataset contains Service Agreement (gas service) level data for random cities. The aggregated data is as of 3/1/2019 for this project's purpose. This will later be updated so that Python will dynamically generate the data being brought in. For now the data is generated from the Queries - Used in Data Wrangling.sql file in this repository. Some code was removed for masking and PII concerns. There are several temp tables created in the mentioned .sql file. At a high level, the data contains:

- Service Agreement level data
- Weather information
- Turn Ons
- Turn Offs
- Credit Scores
- DNP flags
- Number of pledges, payments, bills, late payments, and usage in the past 18 months
- Total money owed
- Prior 3 years of being a seasonal customer
- Prior premise and person level stops/starts.

The data is read into the Python notebook through a csv (data.csv).

I then set the index for the dataset to the unique identifier (SA_ID).

Several fields needed to have their data types updated. I removed Pending Starts, Reactivated, and Cancelled SAs. And also removed a few bad records that had Null Rate Class, Company, Bill Cycle values.

Updated Dataset After Data Wrangling

For the most part, the columns are the same as what was read into the notebook. The big difference is in the data types and removal of some of the records.


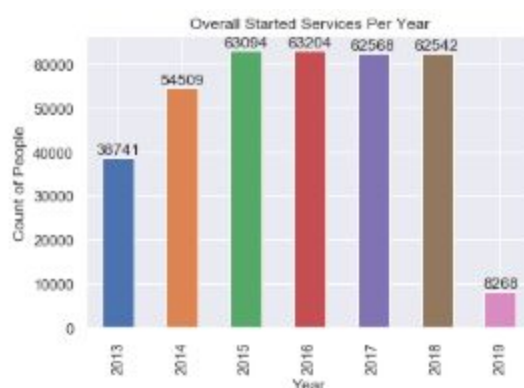
CIS_DIVISION	629614	non-null	object	SA_START_DEGREE_DAY	629614	non-null	int64
ACCOUNT_ID	629614	non-null	object	SA_START_AVG_TEMP	629614	non-null	int64
CUSTOMER_CLASS_CODE	629614	non-null	object	SA_END_DEGREE_DAY	629614	non-null	int64
CUSTOMER_CLASS_DESCRIPTION	629614	non-null	object	SA_END_AVG_TEMP	629614	non-null	int64
PERSON_ID	629614	non-null	object	STOP_2013	629614	non-null	int64
SA_START_DATE	629614	non-null	datetime64[ns]	STOP_2014	629614	non-null	int64
SA_START_YEAR	629614	non-null	object	STOP_2015	629614	non-null	int64
SA_START_YEAR_MONTH	629614	non-null	object	STOP_2016	629614	non-null	int64
SA_END_DATE	629614	non-null	datetime64[ns]	STOP_2017	629614	non-null	int64
SA_END_YEAR	629614	non-null	object	STOP_2018	629614	non-null	int64
SA_END_YEAR_MONTH	629614	non-null	object	STOP_2019	629614	non-null	int64
SA_STATUS_FLAG	629614	non-null	object	START_2013	629614	non-null	int64
SA_TYPE_CODE	629614	non-null	object	START_2014	629614	non-null	int64
RATE_CLASS_CODE	629614	non-null	object	START_2015	629614	non-null	int64
RATE_CLASS_DESCRIPTION	629614	non-null	object	START_2016	629614	non-null	int64
PREMISE_ID	629614	non-null	object	START_2017	629614	non-null	int64
CITY	629614	non-null	object	START_2018	629614	non-null	int64
STATE	629614	non-null	object	START_2019	629614	non-null	int64
POSTAL	629614	non-null	object	PLEDGE_DATE_2013	0	non-null	datetime64[ns]
BILL_CYCLE_CODE	629614	non-null	object	PLEDGE_FLAG_2013	629614	non-null	int64
BILL_CYCLE_DESCRIPTION	629614	non-null	object	PLEDGE_DATE_2014	0	non-null	datetime64[ns]
SERVICE_TYPE_CODE	629614	non-null	object	PLEDGE_FLAG_2014	629614	non-null	int64
PREMISE_TYPE_CODE	629614	non-null	object	PLEDGE_DATE_2015	4423	non-null	datetime64[ns]
TREND_AREA_CODE	629614	non-null	object	PLEDGE_FLAG_2015	629614	non-null	int64
OFFICE_LOCATION	629578	non-null	object	PLEDGE_DATE_2016	5912	non-null	datetime64[ns]
OFFICE_LOCATION_DESCRIPTION	629578	non-null	object	PLEDGE_FLAG_2016	629614	non-null	int64
DNP_STARTED_FLAG	629614	non-null	object	PLEDGE_DATE_2017	5722	non-null	datetime64[ns]
DNP_STOPPED_FLAG	629614	non-null	object	PLEDGE_FLAG_2017	629614	non-null	int64
INTERNAL_CREDIT_RATING	629614	non-null	int64	PLEDGE_DATE_2018	3137	non-null	datetime64[ns]
PREMISE_LEVEL_12_MTH_DNP_FLAG	629614	non-null	object	PLEDGE_FLAG_2018	629614	non-null	int64
ACCOUNT_LEVEL_12_MTH_DNP_FLAG	629614	non-null	object	PLEDGE_DATE_2019	0	non-null	datetime64[ns]
PERSON_RCVD_18_MTHS_PLEDGE	629614	non-null	object	PLEDGE_FLAG_2019	629614	non-null	int64
PAYMENTS_IN_LAST_18_MONTHS	629614	non-null	int64	USAGE_IN_LAST_18_MONTHS	424436	non-null	float64
BILLS_IN_LAST_18_MONTHS	629614	non-null	int64	PERSON_MAX_SA_START_DATE	629614	non-null	datetime64[ns]
PAY_SEGS_IN_LAST_18_MONTHS	629614	non-null	int64	PERSON_MIN_SA_START_DATE	629614	non-null	datetime64[ns]
BILL_SEGS_IN_LAST_18_MONTHS	629614	non-null	int64	PERSON_MAX_SA_END_DATE	368874	non-null	datetime64[ns]
ARREARS_CURRENT_AMOUNT	629614	non-null	float64	PERSON_MIN_SA_END_DATE	368874	non-null	datetime64[ns]
ARREARS_PAYOFF_AMOUNT	629614	non-null	float64	PREMISE_MAX_SA_START_DATE	629614	non-null	datetime64[ns]
TOTAL_CURRENT_AMOUNT	629614	non-null	float64	PREMISE_MIN_SA_START_DATE	629614	non-null	datetime64[ns]
TOTAL_PAYOFF_AMOUNT	629614	non-null	float64	PREMISE_MAX_SA_END_DATE	421190	non-null	datetime64[ns]
MOST_RECENT_PAYMENT_DATE	569558	non-null	datetime64[ns]	PREMISE_MIN_SA_END_DATE	421190	non-null	datetime64[ns]
LATE_PAYMENT_COUNT	629614	non-null	int64	PREMISE_PRIOR_STOP_DATE	259860	non-null	datetime64[ns]
SEASONAL_PRIOR_1_YR_FLAG	629614	non-null	object	PERSON_PRIOR_STOP_DATE	156888	non-null	datetime64[ns]
SEASONAL_PRIOR_2_YR_FLAG	629614	non-null	object	PREMISE_DAYS_INACTIVE_BEFORE	629614	non-null	int64
SEASONAL_PRIOR_3_YR_FLAG	629614	non-null	object	PERSON_DAYS_INACTIVE_BEFORE	629614	non-null	int64
SA_START_DEGREE_DAY	629614	non-null	int64	PREMISE_DAYS_ACTIVE_BEFORE	629614	non-null	int64
				PERSON_DAYS_ACTIVE_BEFORE	629614	non-null	int64
				ACTIVE_DIFF_FROM_20190301	341199	non-null	float64

Field Definitions

- **CIS Division:** Company
- **Account ID:** An account is owned by 1 person but can have multiple Service Agreements.
- **Customer Class:** Account level customer class.
- **Person ID:** A person can have multiple accounts.
- **SA Start:** The Service Agreement's Start Date (when gas started).
- **SA End:** The Service Agreement's End Date (when gas stopped).
- **SA Status:** The current SA status.
- **SA Type:** The Service Agreement's type (example: Gas Residential).
- **Rate Class:** Determines how the SA is billed.
- **Premise ID:** Unique identifier for an address.
- **City:** Premise's City
- **State:** Premise's State
- **Postal:** Premise's Zip Code
- **Bill Cycle:** How we know when to bill an Account.

- **Service Type:** Always Gas.
- **Premise Type:** Premise Type (example: Commercial).
- **Office Location:** Closest office to premise.
- **DNP Started Flag:** Was the previous stop a DNP.
- **DNP Stopped Flag:** Did the current SA Stop because of DNP.
- **Internal Credit Rating:** The internal credit rating for an Account.
- **Premise Level 12 Month DNP Flag:** Has the Premise had a DNP in the last 12 months.
- **Account Level 12 Month DNP Flag:** Has the Account had a DNP in the last 12 months.
- **Person Received 18 Months Pledge:** Has the Person received a Pledge in the last 18 months.
- **Payments in Last 18 Months:** How many payments have been made for the SA in the last 18 months.
- **Bills in Last 18 Months:** How many bills were billed for the SA in the last 18 months.
- **Arrears Current/Payoff Amount:** How much money is past due for the SA.
- **Total Current/Payoff Amount:** How much money is owed for the SA.
- **Most Recent Payment Date:** What was the most recent Payment Date.
- **Late Payment Count:** How many late payments were charged to the SA in the past 18 months.
- **Seasonal Flags (1 year ago, 2 years ago, 3 years ago):** Did the premise turn off between 3/1 and 6/30 and then turn on between 8/1 and 12/31.
- **SA Start/End Degree:** What was the Degree Days on the SA Start and SA End dates.
- **Stop/Start Flags:** Flags to show if the SA started/stopped in a given year.
- **Pledge Date/Flag:** The Pledge Date and Flag for each year.
- **Usage in Last 18 Months:** The Usage billed in the past 18 months.
- **Person/Premise Min/Max SA Start/End Dates:** What are the minimum and maximum start and end dates at a Person or Premise level.
- **Person/Premise Days Active/Inactive Before:** How many days was the previous SA for the Person or Premise Active or Inactive for.
- **Stop Days from 20190301:** Flags to use for training/testing. This should be dynamic but we are hardcoding the value for now.


Overall Started/Stopped Services Per Year



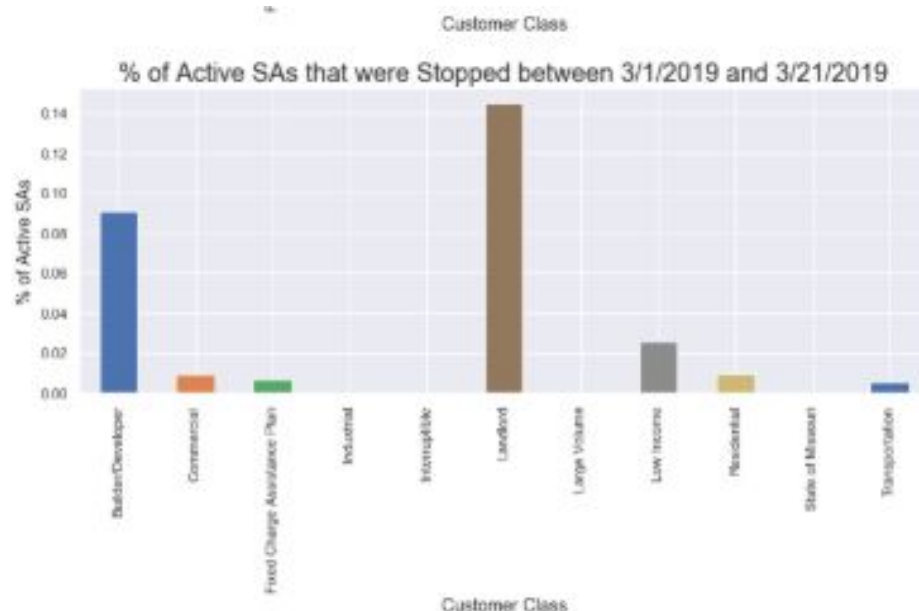
Exploratory Data Analysis

During EDA, several observations of interest were found.


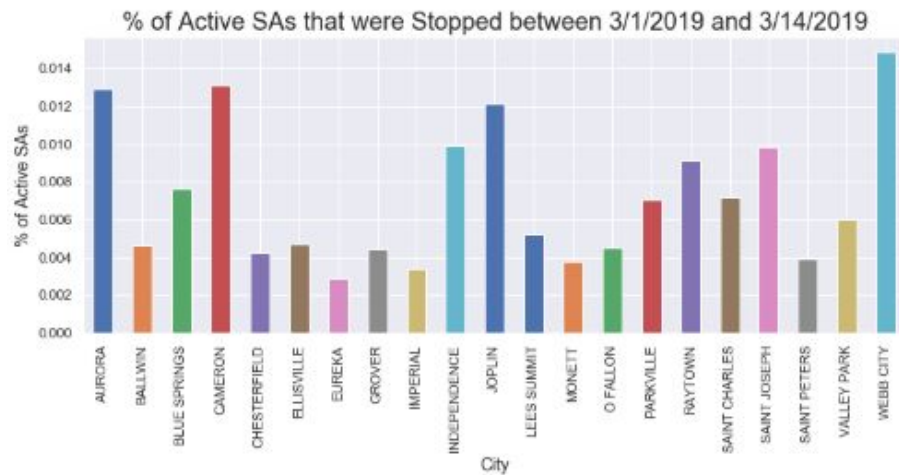
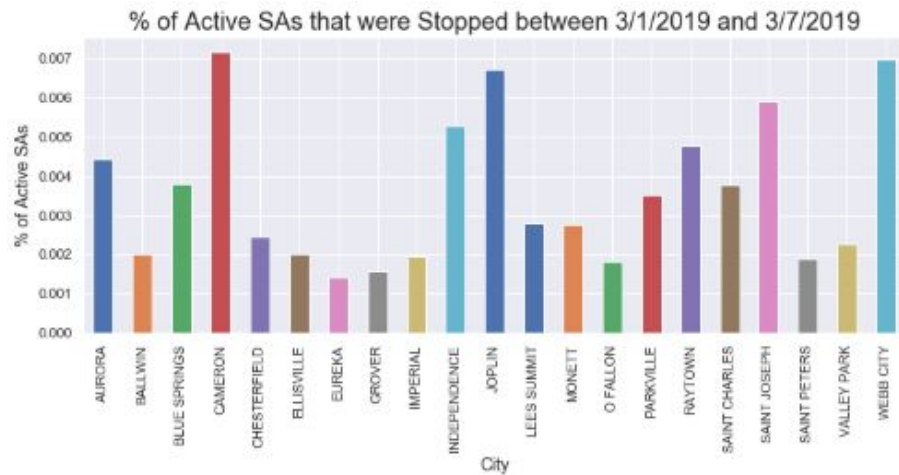
- 1) The most frequent Start Date of the SAs that were stopped in March had started in October, November, and December of the previous year.



- 2) The Residential Customer Class had the most stops during all 3 of the time windows but less than 1% of all active Residential SAs were stopped. The biggest offenders of stops during the 3 windows were Landlords and Builder/Developers. 14% of Landlord SAs were stopped during the 21 day window and 9% of Builder/Developer SAs were stopped during the same window.

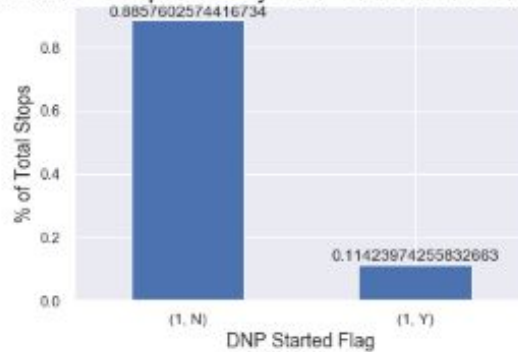


- 3) When looking at stops during the 7 day, 14 day, and 21 day windows by City, Webb City, Joplin, Cameron, and Aurora had the highest % of Active SAs that were stopped. Over 2% of all Active SAs in Aurora were stopped during the 21 day window.




- 4) The percentage of stopped SAs that were DNPs previously during the 7 day window was 11%, 14 day window was 15%, and 21 day window was also 15%.


% of Stopped SAs that were previously DNP or not between 3/1/2019 and 3/7/2019



% of Stopped SAs that were previously DNP or not between 3/1/2019 and 3/14/2019



% of Stopped SAs that were previously DNP or not between 3/1/2019 and 3/21/2019



Statistical Inference

To confirm some of the above observations that may prove to be useful, I applied inferential statistics and found the correlations that are of statistical significance and a few that have practical significance. Here are a few observations:

- When comparing the Pearson correlation between Late Payment Counts correlated to Stopping 21 Days from 3/1/2019 using the full dataset, there is statistical significance in the correlations and it shows that there is little practical significance since r^2 is small ($< .1\%$). But when using the dataset with only Active SAs as of 3/1/2019, the p-value is .13 which indicates there is no significant correlation.
- # of Days Active as of 3/1/2019 have a negative correlation with Stopping. So in other words, the less days you have been active, the more likely you are to stop with 21 days of 3/1/2019.

	ACTIVE_DIFF_FROM_20190301	STOP_7_DAYS_FROM_20190301	STOP_14_DAYS_FROM_20190301	STOP_21_DAYS_FROM_20190301
ACTIVE_DIFF_FROM_20190301	1.000000	-0.039183	-0.052372	-0.062117
STOP_7_DAYS_FROM_20190301	-0.039183	1.000000	0.719951	0.598378
STOP_14_DAYS_FROM_20190301	-0.052372	0.719951	1.000000	0.831137
STOP_21_DAYS_FROM_20190301	-0.062117	0.598378	0.831137	1.000000

- The Bootstrap Estimate of SA Stops Mean shows the 95% confidence interval for SAs stopping between 3/1/2017 and 3/7/2019 is from .34% to .38%. Similarly, 3/1/2017 to 3/14/2019 is from .67% to .72% and 3/1/2019 to 3/21/2019 is from .98% to 1.04%.

